

Mean training duration per algorithm — one panel per pair (averaged across scenarios, n_unknown, seeds)

